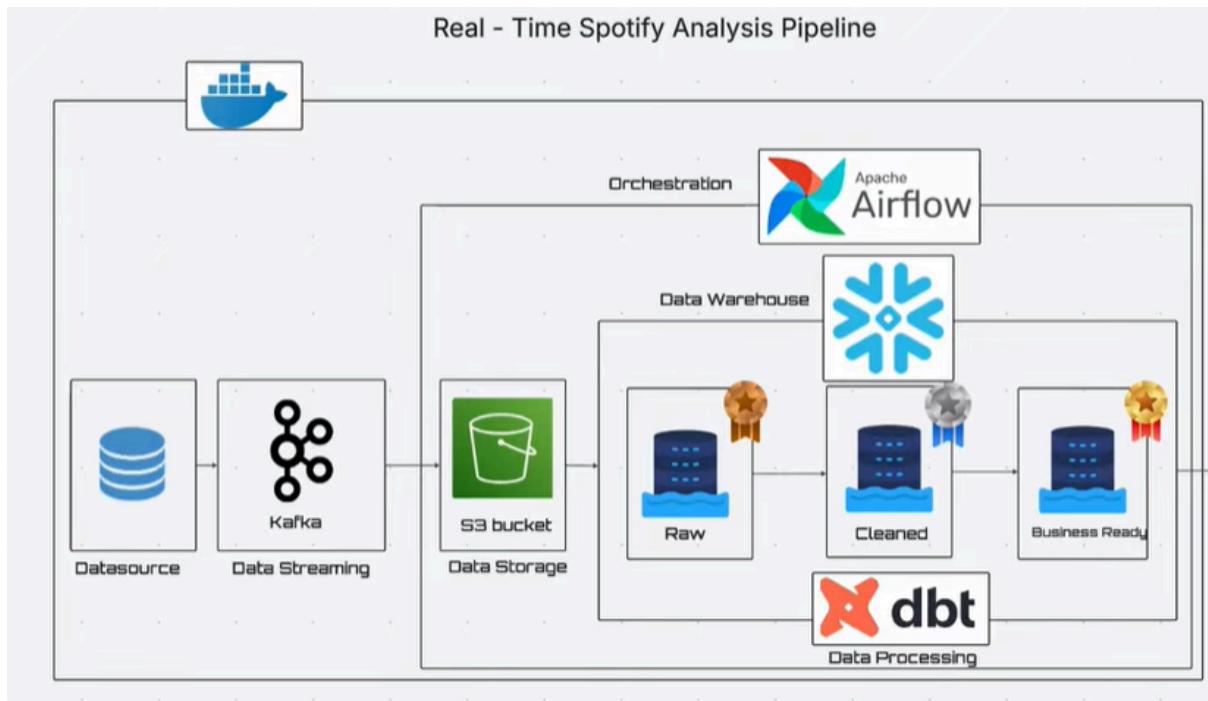


# SPOTIFY DATA PIPELINE

Architecture:



Tech-Stack Used:

- 1.Kafka
- 2.Snowflake
- 3.Airflow
- 4.Dbt
- 5.Docker
- 6.AWS S3

## Development

1. Create a dir simulator (mkdir simulator), this will contain the producer code that will produce data in real-time.
- 2.Create docker dir which will have a docker-compose file having all the containers: zookeeper, kafka, kafdrop(for UI), minio for storage, postgres for initializing airflow & airflow containers. We will also create a .env file.(Make your credentials)
3. Create the requirement.txt file and install all dependencies:  
`pip install -r requirements.txt # here -r is for read`
4. Do docker compose up after changing dir to docker.
5. Check if kafdrop, airflow and minio are working by clicking the port in docker desktop. Check login creds as well.
6. Now let's come to producer code:

6.1 We set up Kafka config and created a Kafka Producer with those config and a value serializer that will convert the dictionary to json.

6.2 Then we created a song\_artist\_pair list of dict, & for each artist-song pair we generate a uuid(universally unique identifier) [uuid5: same song\_id each time]. A **namespace** is just a fixed UUID used as a **seed** to generate other UUIDs in a consistent way.

Key properties of **uuid5**:

- **Deterministic** (same input → same UUID every time)
- **Not random**
- Ideal for **natural keys** → **surrogate IDs**

uuid4 is used to create “random” uuid

6.3 Then we generated random users, and created a “generate\_event” function to generate random records using the above data.

7. **Kafka Consumer**: We will first connect to s3 service and check if a bucket already exists, else we create a new one.

7.1 Initialize a Kafka Consumer and store the events in s3 bucket according to batch size.

7.2 auto\_offset\_reset="earliest" means it will continue from where it last stopped

7.3 kafka\_group\_id - It saves the offset , if we change this whole offset will start again

7.4 json.load will convert json to dict, take a string as input and returns a dictionary as output.

7.5 json.dumps will convert dict to json, take a dictionary as input and returns a string as output.

```
Produced event: pause - God's Plan by Drake (user 881dc191-47cf-4642-b4b2-9334c488cd67)
Produced event: skip - Love Story by Taylor Swift (user dc9d0fc7-9cba-480b-a0d8-093db735f4e1)
Produced event: add_to_playlist - God's Plan by Drake (user 881dc191-47cf-4642-b4b2-9334c488cd67)
Produced event: pause - Stronger by Kanye West (user 384d94c7-154e-4602-8d08-8c665b3de7b9)
Produced event: skip - Levitating by Dua Lipa (user dc9d0fc7-9cba-480b-a0d8-093db735f4e1)
Produced event: add_to_playlist - Love Story by Taylor Swift (user f78657c8-14a0-46b8-9443-04f5d3a411ff)
Produced event: pause - Stronger by Kanye West (user 881dc191-47cf-4642-b4b2-9334c488cd67)
Produced event: pause - Stronger by Kanye West (user fd69b2ea-cc02-4d21-b918-fe873c4db470)
Produced event: skip - Levitating by Dua Lipa (user 19d48e3d-20fc-45e7-8454-04fc8b10a7c1)
Produced event: skip - Love Story by Taylor Swift (user c0d8a451-a49a-41d6-b5fd-8235a6a2e18d)
Produced event: play - Stronger by Kanye West (user 71ffffa9f-1423-4495-a311-74e985682796)
Produced event: play - Stronger by Kanye West (user c0d8a451-a49a-41d6-b5fd-8235a6a2e18d)
Produced event: skip - Levitating by Dua Lipa (user 19d48e3d-20fc-45e7-8454-04fc8b10a7c1)
[]
```

producer v1\*

```

Bucket spotify already exists.
Listening of events on Kafka topic 'spotify-events'.....
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-52-36.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-57-43.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-57-53.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-03.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-14.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-24.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-34.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-44.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-58-54.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-59-04.json
Uploaded 10 events to MinIO: bronze/date=2026-02-15/hour=17/spotify_events_2026-02-15T17-59-14.json

```

Object Browser

spotify

Created on: Sun, Feb 15 2026 23:19:45 (GMT+5:30) Access: PRIVATE 40.7 KIB - 13 Objects

Name	Last Modified	Size
spotify_events_2026-02-15T17-52-36.json	Today, 23:22	3.2 KIB
spotify_events_2026-02-15T17-57-43.json	Today, 23:27	3.1 KIB
spotify_events_2026-02-15T17-57-53.json	Today, 23:27	3.1 KIB
spotify_events_2026-02-15T17-58-03.json	Today, 23:28	3.1 KIB
spotify_events_2026-02-15T17-58-14.json	Today, 23:28	3.2 KIB
spotify_events_2026-02-15T17-58-24.json	Today, 23:28	3.2 KIB
spotify_events_2026-02-15T17-58-34.json	Today, 23:28	3.2 KIB
spotify_events_2026-02-15T17-58-44.json	Today, 23:28	3.1 KIB
spotify_events_2026-02-15T17-58-54.json	Today, 23:28	3.1 KIB
spotify_events_2026-02-15T17-59-04.json	Today, 23:28	3.1 KIB
spotify_events_2026-02-15T17-59-14.json	Today, 23:28	3.1 KIB

Now we have data in s3 bucket

## 8. Move data from S3 to Snowflake using Airflow:

8.1 Create a DAG (spotify\_minio\_to\_snowflake\_bronze)

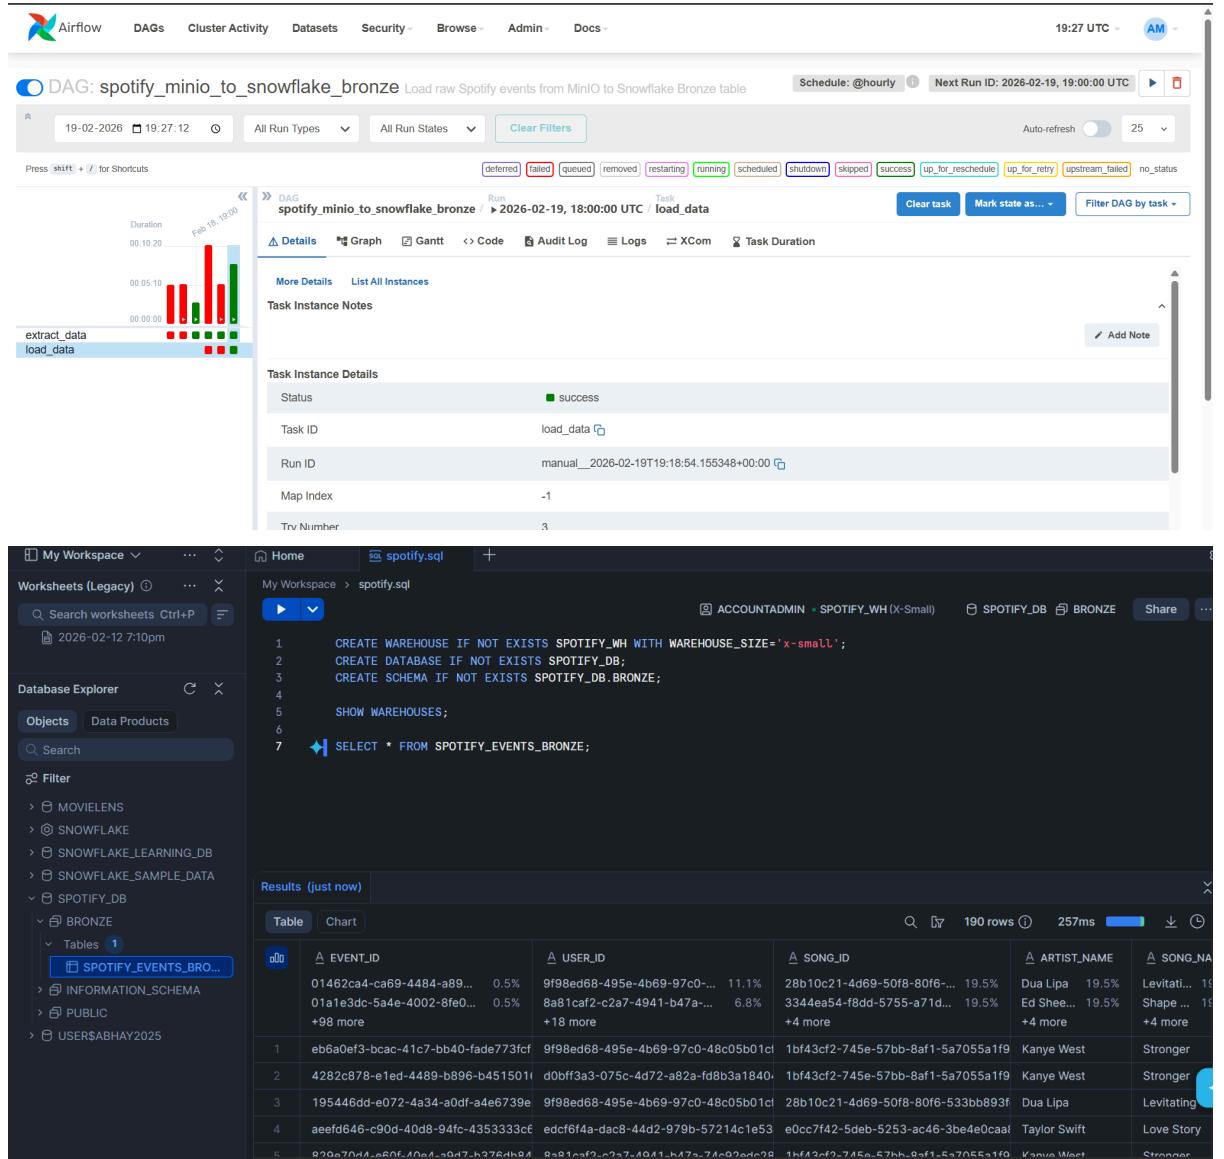
8.2 Two tasks : extract\_task to extract data from s3, and load\_task to load data in snowflake bronze layer.

8.3 Extract task will extract json files data and store it into temp file on local path

8.4 load\_to\_snowflake will load raw data in the SF bronze layer for further processing.

\*\*context : Airflow context dictionary that contains metadata and information about a running Directed Acyclic Graph (DAG), its tasks, and the Airflow environment.

Below is the successful run of Airflow DAG which loaded data into snowflake DB bronze layer table:



## 9. DBT:

### 9.1 Create a new folder spotify\_dbt

9.2 Initialize dbt: dbt init and enter the details of snowflake account , warehouse, database, default schema(we used Transform schema created as default), threads, it will create a profile.

9.3 Create sources.yml file containing metadata for silver layer source.

9.4 Create silver dir & spotify\_sql table to transform data from bronze layer and store in silver layer.

9.5 Similarly create models for the gold layer.

9.6 Run >> dbt run inside spotify\_dbt dir

```

26 -- Create new schema for bronze layer
27 CREATE SCHEMA IF NOT EXISTS SPOTIFY_DB.MART;
28 USE SCHEMA MART;
29
30 ┆ SELECT * FROM MART.TOP_SONGS;
31

```

	SONG_ID	SONG_NAME	ARTIST_NAME	TOTAL_PLAYS	TOTAL_SKIPS
1	3344ea54-f8dd-5755-a71d-2083035fb139	Shape of You	Ed Sheeran	74	74
2	28b10c21-4d69-50f8-80f6-533bb893fe01	Levitating	Dua Lipa	74	74
3	1fb43cf2-745e-57bb-8af1-5a7055a1f93a	Stronger	Kanye West	66	66
4	09a7d5ca-d131-55b1-9758-8cfcd6375651	God's Plan	Drake	66	66
5	e0cc7f42-5deb-5253-ac46-3be4e0caa845	Love Story	Taylor Swift	56	56
6	74a42680-ff06-55cf-9f7b-426cd71bda2b	Blinding Lights	The Weeknd	44	44

```

26 -- Create new schema for bronze layer
27 CREATE SCHEMA IF NOT EXISTS SPOTIFY_DB.MART;
28 USE SCHEMA MART;
29
30 ┆ SELECT * FROM MART.USER_ENGAGEMENT;
31

```

	USER_ID	DEVICE_TYPE	COUNTRY	TOTAL_PLAYS	TOTAL_SKIPS	PLAYLIST_ADDS	DAY
1	8ab1caf2-c2a7-4941-b47a-74c9... 7.5%	mobile	CA 20.5%	34.9%	32.9%	10	10
2	9f98ed68-495e-4b69-97c0-48c05b01cf67	desktop	DE 19.2%	+1 more	+4 more	2	2
3	9f98ed68-495e-4b69-97c0-48c05b01cf67	web	CA	6	6	6	6
4	7bcd17fe-29bc-4ef5-b730-212c2922c387	mobile	IN	6	6	6	6
5	c162de85-8cfe-40c9-889e-69521651f727	web	UK	6	6	6	6
6	65daca3d-cfd6-48f1-960b-6764aee3d3d	desktop	CA	4	4	4	4
7	9f98ed68-495e-4b69-97c0-48c05b01cf67	web	IN	4	4	4	4
8	c7d7fd5a-0ddc-4ab1-a0e8-8f28973f3df5	web	DE	4	4	4	4
9	c319e0da-5063-4d95-80c0-e597e6006317	mobile	DE	4	4	4	4

Note: Please make sure to update .env files with correct username, password and account of snowflake before running.

Thanks for reading the documentation.